

Essential Linux Command Reference Guide

Command 1: uname

(1) What is the uname command? (History, invention, author)

The uname command is a terminal utility program used in Unix/Linux to display detailed information about the system including kernel name, version, hardware platform, and operating system details.

History & Origin

- The uname command first appeared in **PWB/UNIX (Programmer's Workbench UNIX)**
- Both the system call and command are specified by **POSIX (Portable Operating System Interface)** standards
- **Author:** David MacKenzie wrote and developed the GNU version of uname
- The GNU version is included in the "coreutils" or "sh-utils" packages
- **Current Status:** uname is not available as a standalone program; it comes bundled with coreutils
- **Available across platforms:** Linux, macOS, Windows (via GnuWin32 project and UnxUtils)

Purpose

To allow system administrators and users to:

- Identify the operating system and kernel details
- Check system architecture and processor type
- Display hardware platform information
- Gather system metadata for scripting and automation
- Verify system compatibility before running specific software

(2) Why is uname useful for System Administrators?

System administrators frequently work with:

- Multiple Linux systems with different configurations
- Automated deployment and configuration scripts
- Software compatibility verification
- Cross-platform system environments
- Automated infrastructure management

Reasons for Use

- **Quick System Identification:** Shows OS name, kernel version, and hardware details instantly without complex commands
- **Scripting Support:** Output can be parsed in shell scripts for conditional execution based on system type

- **Compatibility Checking:** Helps verify if a system meets software requirements before installation
- **Multi-system Management:** Essential for managing heterogeneous environments (Linux, macOS, BSD)

(3) Basic Syntax and Examples

Basic Syntax:

```
```bash
uname [OPTIONS]
```

```

Common Options:

| Option | Description |
|--------|---|
| (none) | Prints kernel name (default) |
| -a | All information (kernel, hostname, kernel version, processor, OS) |
| -s | Kernel name |
| -n | Network node hostname |
| -r | Kernel release |
| -v | Kernel version |
| -m | Machine hardware name |
| -p | Processor type |
| -i | Hardware platform |
| -o | Operating system |

Table 1: uname command options

Example 1: Display all system information

```
```bash
$ uname -a
Linux localhost.localdomain 5.10.0-14-generic #1 SMP Fri Oct 22 09:03:12 UTC 2021 x86_64
x86_64 x86_64 GNU/Linux
```

```

Example 2: Check kernel release for compatibility

```
```bash
$ uname -r
5.10.0-14-generic
```

```

Example 3: Display processor architecture

```
```bash
$ uname -m
x86_64
````
```

Example 4: Scripting with uname

```
```bash
#!/bin/bash
if [[$(uname -s) == "Linux"]]; then
echo "Running on Linux system"
KERNEL=$(uname -r)
echo "Kernel version: $KERNEL"
else
echo "Not running on Linux"
fi
````
```

(4) Practical Use Cases

- **Verify system before software installation:** Check architecture compatibility
- **Monitor fleet of servers:** Quickly identify OS versions across infrastructure
- **Automate system detection:** Write portable shell scripts that adapt to the running OS
- **Troubleshooting:** Identify kernel or hardware issues

Command 2: pgrep

(1) What is the pgrep command? (History, invention, author)

The pgrep command is a terminal utility program used in Unix/Linux to search for running processes based on specified criteria (name, user ID, terminal, etc.) and display their process IDs (PIDs) to standard output.

History & Origin

- pgrep is part of the **procps-ng package** (successor to procps package)
- Also available in **util-linux** and **procutil** packages depending on Linux distribution
- Originally developed as a more efficient alternative to parsing ps output with grep
- First appeared in **BSD Unix**, later adopted by Linux distributions worldwide
- Primary maintainers: Linux community and BSD developers
- Highly optimized for process searching without loading entire process list
- Available on Linux, BSD, macOS, and other Unix-like systems

Purpose

To allow users and administrators to:

- Find process IDs by process name without complex piping commands
- Search processes by user, terminal, or other attributes efficiently
- Count running instances of a specific process
- Kill specific process instances more safely than using wildcards
- Automate process monitoring and management in scripts

- Create condition-based scripts that check for process existence

(2) Why is pgrep useful for System Administrators?

System administrators frequently work with:

- Running services and application monitoring
- Process checking in automated cron jobs
- Service status verification for critical applications
- Process termination and cleanup tasks
- Application debugging and troubleshooting
- Resource allocation based on running processes

Reasons for Use

- **Simpler Alternative to ps | grep:** Much cleaner syntax compared to complex pipe operations
- **Avoids Matching grep Itself:** Prevents false matches that occur when using grep to search for processes
- **Direct PID Output:** Returns only process IDs, not entire process lines

(3) Basic Syntax and Examples

Basic Syntax:

```
```bash
pgrep [OPTIONS] PATTERN
```

```

Common Options:

| Option | Description |
|-------------|---|
| -l | Show process name along with PID |
| -u USER | Search only for processes owned by USER |
| -U USER | Search only for processes NOT owned by USER |
| -t TERMINAL | Search only on specified terminal |
| -c | Count matching processes |
| -f | Match against full process command line |
| -g PID | Match process group ID |
| -P PID | Match parent process ID |
| -x | Match exact process name |
| -n | Return newest matching process |

Table 2: pgrep command options

Example 1: Find all processes named 'apache2'

```
```bash
$ pgrep apache2
1234
1235
1236
```

```

Example 2: Find processes with names and PIDs

```
```bash
$ pgrep -l apache2
1234 apache2
1235 apache2
1236 apache2
```

```

Example 3: Count running processes of a specific user

```
```bash
$ pgrep -u pranav -c
23
```

```

Example 4: Find newest instance of a process

```
```bash
$ pgrep -n python3
5678
```

```

Example 5: Kill all processes matching a pattern

```
```bash
$ pkill -f "old_script.py"
```

```

(4) Practical Use Cases

- **Monitor service availability:** Check if critical processes are running
- **Automated startup scripts:** Prevent duplicate processes from starting
- **System cleanup:** Terminate zombie or hung processes efficiently
- **Log monitoring:** Send alerts if important services stop running

Command 3: ps (Process Status)

What is ps?

The **ps** command displays information about active processes currently running on the system. It provides detailed snapshots of process status including memory usage, CPU usage, and execution status.

Syntax:

```
```bash
ps [OPTIONS]
```
```

Common Use Cases:

- View running processes with detailed information
- Monitor resource consumption (CPU, memory)
- Track process hierarchy and parent-child relationships
- Identify zombie processes
- Track user-specific processes

Examples:

```
```bash
```

## Show all processes

```
$ ps aux
```

## Show process tree

```
$ ps -ef
```

## Show processes with memory usage sorted

```
$ ps aux --sort=%mem
```

## Show specific user's processes

```
$ ps -u username
```

## Show process hierarchy

```
$ ps --forest -ef
```
```

Command 4: top (Real-time Process Monitor)

What is top?

The **top** command provides real-time monitoring of system processes, showing CPU and memory usage, load averages, and process details in an interactive interface.

Key Features:

- Real-time process monitoring
- Interactive control (pause, resume, sort)
- System load and memory statistics

- Customizable display and sorting
- Process priority adjustment

Interactive Commands:

- **q:** Quit top
- **k:** Kill a process (specify PID)
- **r:** Change process priority (renice)
- **f:** Add/remove columns
- **P:** Sort by CPU usage
- **M:** Sort by memory usage

Example Usage:

```
```bash
```

## **Launch top monitor**

```
$ top
```

## **Launch top and exit after 2 iterations**

```
$ top -n 2
```

## **Monitor specific process by PID**

```
$ top -p 1234
```

## **Batch mode output**

```
$ top -b -n 1
```

```

Command 5: kill and pkill (Process Termination)

Terminating Processes Safely

The **kill** command sends signals to processes. The most common signals are:

| Signal | Name | Effect |
|--------|---------|---------------------------------|
| 1 | SIGHUP | Hang up / restart |
| 2 | SIGINT | Interrupt (Ctrl+C) |
| 9 | SIGKILL | Force kill (cannot be caught) |
| 15 | SIGTERM | Terminate gracefully (default) |
| 19 | SIGSTOP | Stop process (cannot be caught) |

Table 3: Common kill signals

Safe Termination Process:

1. First attempt graceful termination (SIGTERM, signal 15)
2. Wait a few seconds for process cleanup
3. If necessary, force termination (SIGKILL, signal 9)

Examples:

```
```bash
```

## Graceful termination

```
$ kill -15 1234
```

## Force termination

```
$ kill -9 1234
```

## Kill by process name

```
$ pkill -f "python script.py"
```

## Kill all processes of a user

```
$ pkill -u username
```

```
```
```

Command 6: systemctl (System Service Management)

Managing System Services

The **systemctl** command controls the systemd system and service manager, essential for modern Linux distributions (Ubuntu, CentOS, Fedora, etc.).

Common Operations:

| Command | Purpose |
|---------------------------|----------------------|
| systemctl start SERVICE | Start a service |
| systemctl stop SERVICE | Stop a service |
| systemctl restart SERVICE | Restart a service |
| systemctl reload SERVICE | Reload configuration |
| systemctl enable SERVICE | Enable on boot |
| systemctl disable SERVICE | Disable on boot |
| systemctl status SERVICE | Check service status |
| systemctl list-units | List all services |

Table 4: systemctl commands

Examples:

```
``bash
```

Check if Apache is running

```
$ sudo systemctl status apache2
```

Start the web server

```
$ sudo systemctl start apache2
```

Enable Apache to start on boot

```
$ sudo systemctl enable apache2
```

Restart PostgreSQL database

```
$ sudo systemctl restart postgresql
```

View all running services

```
$ systemctl list-units --type=service --state=running
``
```

Real-World Scenarios and Scripts

Scenario 1: Monitoring Web Server Health

```
```bash
#!/bin/bash

SERVICE="apache2"
ADMIN_EMAIL="admin@example.com"
```

## Check if service is running

```
if pgrep -x "$SERVICE" > /dev/null; then
 echo "$SERVICE is running"
else
 echo "$SERVICE is not running - attempting restart"
 sudo systemctl restart $SERVICE
```

```
Check again
if pgrep -x "\$SERVICE" > /dev/null; then
 echo "\$SERVICE restarted successfully"
else
 echo "ALERT: \$SERVICE failed to start"
 # Send alert email
fi
```

```
fi
```

```

Scenario 2: System Resource Audit

```
```bash
#!/bin/bash

echo "==== System Information ===="
uname -a
echo ""

echo "==== Top Memory Consumers ===="
ps aux --sort=-%mem | head -6
echo ""

echo "==== Top CPU Consumers ===="
ps aux --sort=-%cpu | head -6
echo ""

echo "==== Process Count by User ===="
pgrep -u . -c | sort -rn | head -5
```

...

### Scenario 3: Automated Cleanup Script

```
```bash
#!/bin/bash
```

Kill processes older than 7 days without activity

```
OLD_PROCESSES=$(pgrep -f "inactive_task")
```

```
for PID in $OLD_PROCESSES; do
echo "Terminating old process: $PID"
kill -15 $PID
```

```
# Force kill if not terminated
sleep 2
if kill -0 \$PID 2>/dev/null; then
    kill -9 \$PID
fi
```

```
done
```

```

## Best Practices and Tips

- **Always use graceful termination first:** Use SIGTERM (signal 15) before SIGKILL (signal 9)
- **Verify processes before killing:** Double-check the PID and process name
- **Use pgrep for scripts:** More reliable than ps | grep for automation
- **Monitor systemd services:** Use systemctl status regularly
- **Document critical processes:** Know which services are essential for your system
- **Test in non-production:** Always test process management scripts in development first
- **Keep logs:** Monitor syslog for process start/stop events
- **Use process monitoring tools:** Consider tools like monit, nagios for production systems

---

## Conclusion

Understanding these essential Linux commands—**uname**, **pgrep**, **ps**, **top**, **kill**, and **systemctl**—is crucial for effective system administration. Each command serves a specific purpose in process management and system monitoring. Combined, they provide administrators with comprehensive tools to manage, monitor, and troubleshoot system processes efficiently.

Mastering these commands enables you to:

- Quickly diagnose system issues
- Automate routine administrative tasks
- Maintain system stability and performance
- Respond promptly to service failures
- Optimize resource utilization

Regular practice with these tools will improve your Linux administration skills and system management capabilities.

## References

- [1] Linux man-pages online documentation: <https://man7.org/linux/man-pages/>
- [2] GNU Coreutils - uname documentation: <https://www.gnu.org/software/coreutils/manual/>
- [3] procps-ng package documentation: <https://gitlab.com/procps-ng/procps>
- [4] systemd system and service manager: <https://systemd.io/>
- [5] Advanced Bash-Scripting Guide: Process Management chapter: <http://tldp.org/LDP/abs/html/process-sub.html>
- [6] Red Hat System Administration Guide: Managing Processes: <https://access.redhat.com/documentation/>